

MPI Version of the Stommel Code with One and Two Dimensional Decomposition

Timothy H. Kaiser, Ph.D.

tkaiser@sdsc.edu

Overview

We will choose one of the two dimensions and subdivide the domain to allow the distribution of the work across a group of distributed memory processors

We will focus on the principles and techniques used to do the MPI work in the model

STEP1: introduce the MPI environment

Need to include “mpif.h” to define MPI constants

Need to define our own constants

numnodes - how many processors are running

myid - Which processor am I

mpi_err - error code returned by most calls

mpi_master- the id for the master node

Suggestion - add the following module to your source and “use” it in the program stommel

```
module mpi
  include "mpif.h"
  integer numnodes, myid, mpi_err
  integer, parameter::mpi_master=0
end module
```

STEP1: Start the MPI environment

Suggestion - add the following to your program

```
call MPI_INIT( mpi_err )  
call MPI_COMM_SIZE(MPI_COMM_WORLD,numnodes,  
mpi_err)  
call MPI_COMM_RANK(MPI_COMM_WORLD, myid,  
mpi_err)  
write(*,*)"from ", myid,"numnodes=",numnodes
```

To stop add the following next

```
call MPI_Finalize(mpi_err)  
stop
```

STEP2: Try it!

Compile

```
f90 -ffree stommel.f -o st_1d
```

Run

```
mpprun -n 3 st_1d
```

Try running again. Do you get the same output?

Input

We read the data on processor 0 and send to the others

```
if(myid .eq. mpi_master)then
    read(*,*)nx,ny
    read(*,*)lx,ly
    read(*,*)alpha,beta,gamma
    read(*,*)steps
endif
```

We use MPI_BCAST to send the data to the other processors

We use 8 calls

Can you do it in 2?


How about 1?

Domain Decomposition (1d)

Physical domain is sliced into sets of columns so that computation in each set of columns will be handled by different processors. Why do columns and not rows?

Serial Version

all cells on one processor




Parallel Version

node 0

node 1

node 2



Domain Decomposition (1d)

Fortran 90 allows us to set arbitrary bounds on arrays

We set our array bounds differently on each processor so that:

We take our original grid and break it into numnodes subsections of size $nx/numnodes$

Each processor calculates for a different subsection of the grid

No two processors calculate ψ for the same (I,J)


For whichever processor hold $\Psi(I,J)$ it corresponds exactly to $\Psi(I,J)$ in the serial program

We add special boundary cells for each subsection of the grid called ghost cells

The values for the ghost cells are calculated on neighboring processors and sent using MPI calls.

Domain Decomposition (1d)

With ghost cells out decomposition becomes...




Domain Decomposition (1d)

How and why are ghost cells used?

Node 0 allocates space for $\text{psi}(0:7,0:3)$ but calculates $\text{psi}(1:6,1,2)$

Node 1 allocates space for $\text{psi}(0:7,2:5)$ but calculates $\text{psi}(1:6,3,4)$

Node 2 allocates space for $\text{psi}(0:7,4:7)$ but calculates $\text{psi}(1:6,5,6)$



To calculate the value for $\text{psi}(4,4)$ node1 requires the value from $\text{psi}(4,3), \text{psi}(5,4), \text{psi}(3,4), \text{psi}(4,5)$

Where does it get the value for $\text{psi}(4,5)$? From node2 and it holds the value in a ghost cell

Domain Decomposition (1d)

Source code for setting up the distributed grid with ghost cells

```
! we stripe the grid across the processors

i1=1
i2=ny
dj=real(nx,b8)/real(numnodes,b8)
j1=nint(1.0_b8+myid*dj)
j2=nint(1.0_b8+(myid+1)*dj)-1
write(*,101)myid,i1,i2,j1,j2
101 format("myid= ",i3,3x,          &
          " (",i3," <= i <= ",i3,") , ", &
          " (",i3," <= j <= ",i3,")")

! allocate the grid to (i1-1:i2+1,j1-1:j2+1) this includes boundary
cells

allocate(psi(i1-1:i2+1,j1-1:j2+1))
```

Try adding this to your program. What do you get?

Ghost cell updates

When do we update ghost cells?

Each trip through our main loop we call `do_transfer` to update the ghost cells

Our main loop becomes...

```
do i=1,steps
  call do_jacobi(psi,new_psi,mydiff,i1,i2,j1,j2)
  call do_transfer(psi,i1,i2,j1,j2)
  write(*,*)i,diff
enddo
```

How do we update ghost cells?

**Processors send and receive values to and from neighbors
Need to exchange with left and right neighbors except
processors on far left and right only transfer in 1 direction**

Trick 1 to avoid deadlock

Even # processors

send left
receive from left
send right
receive from right

Odd # processors

receive from right
send to right
receive for left
send to left

Trick 2 to handle the end processors

Send to MPI_PROC_NULL instead of a real processor

How do we update ghost cells?

! How many cells are we sending

```
num_x=i2-i1+3
```

! Where are we sending them

```
myleft=myid-1
```

```
myright=myid+1
```

```
if(myleft .le. -1)myleft=MPI_PROC_NULL
```

```
if(myright .ge. numnodes)myright=MPI_PROC_NULL
```

How do we update ghost cells?

For even numbered processors...

```
if(even(myid))then
! send to left
  call MPI_SEND(psi(:,j1), num_x,MPI_DOUBLE_PRECISION,myleft, &
    100,MPI_COMM_WORLD,mpi_err)
! rec from left
  call MPI_RECV(psi(:,j1-1),num_x,MPI_DOUBLE_PRECISION,myleft, &
    100,MPI_COMM_WORLD,status,mpi_err)
! rec from right
  call MPI_RECV(psi(:,j2+1),num_x,MPI_DOUBLE_PRECISION,myright, &
    100,MPI_COMM_WORLD,status,mpi_err)
! send to right
  call MPI_SEND(psi(:,j2), num_x,MPI_DOUBLE_PRECISION,myright, &
    100,MPI_COMM_WORLD,mpi_err)
else
```

How do we update ghost cells?

For odd numbered processors...

```
Else ! we are on an odd column processor
! rec from right
  call MPI_RECV(psi(:,j2+1),num_x,MPI_DOUBLE_PRECISION,myright, &
               100,MPI_COMM_WORLD,status,mpi_err)
! send to right
  call MPI_SEND(psi(:,j2), num_x,MPI_DOUBLE_PRECISION,myright, &
               100,MPI_COMM_WORLD,mpi_err)
! send to left
  call MPI_SEND(psi(:,j1), num_x,MPI_DOUBLE_PRECISION,myleft, &
               100,MPI_COMM_WORLD,mpi_err)
! rec from left
  call MPI_RECV(psi(:,j1-1),num_x,MPI_DOUBLE_PRECISION,myleft, &
               100,MPI_COMM_WORLD,status,mpi_err)
endif
```


How do we update ghost cells?

It's a 4 stage operation

Example with 4 nodes:

	Node 0	Node 1	Node 2	Node 3
Stage 1	Send left to MPI_PROC_NULL	Receive right from 2	Send left to 2	Receive right from MPI_PROC_NULL
Stage 2	Receive left from MPI_PROC_NULL	Send right to 2	Receive left from 2	Send right to MPI_PROC_NULL
Stage 3	Receive right from 1	Send left to 0	Receive right from 3	Send left to 0
Stage 4	Send left to 1	Receive left from 0	Send left to 3	Receive left from 0

Only a few other modifications

Force and do_jacobi are not modified

We modify the boundary condition routine to only set value for true boundaries and ignore ghost cells

```
subroutine bc(psi,i1,i2,j1,j2)
! sets the boundary conditions
! input is the grid and the indices for the interior cells
  use numz
  use mpi
  use input, only : nx,ny
  implicit none
  real(b8),dimension(i1-1:i2+1,j1-1:j2+1):: psi
  integer,intent(in):: i1,i2,j1,j2
! do the top edges
  if(i1 .eq. 1) psi(i1-1,:)=0.0_b8
! do the bottom edges
  if(i2 .eq. ny) psi(i2+1,:)=0.0_b8
! do left edges
  if(j1 .eq. 1) psi(:,j1-1)=0.0_b8
! do right edges
  if(j2 .eq. nx) psi(:,j2+1)=0.0_b8
end subroutine bc
```

Residual

In our serial program the routine `do_jacobi` calculates a residual for each iteration

The residual is the sum of changes to the grid for an jacobi iteration

Now the calculation is spread across all processors

To get the global residual we can use the `MPI_Reduce` function

```
call MPI_REDUCE(mydiff,diff,1,MPI_DOUBLE_PRECISION, &  
               MPI_SUM,mpi_master,MPI_COMM_WORLD,mpi_err)  
if(myid .eq. mpi_master)write(*,*)i,diff
```

Our main loop is now...

Call the `do_jacobi` subroutine

Update the ghost cells

Calculate the global residual

```
do i=1,steps
  call do_jacobi(psi,new_psi,mydiff,i1,i2,j1,j2)
  call do_transfer(psi,i1,i2,j1,j2)
  call MPI_REDUCE(mydiff,diff,1,MPI_DOUBLE_PRECISION, &
                 MPI_SUM,mpi_master,MPI_COMM_WORLD,mpi_err)
  if(myid .eq. mpi_master)write(*,*)i,diff
enddo
```

Final change

We change the `write_grid` subroutine so that each node writes its part of the grid to a different file.

Function `unique` returns a file name based on a input string and the node number

We change the open statement in `write_grid` to:

```
open(18,file=unique("out1d_"),  
     recl=max(80,15*((jend-jstart)+3)+2))
```

Unique

We add an interface to unique in the module face
Unique is the function:

```
function unique(name)
  use numz
  use mpi
  character (len=*) name
  character (len=20) unique
  character (len=80) temp
  if(myid .gt. 99)then
    write(temp,"(a,i3)")trim(name),myid
  else
    if(myid .gt. 9)then
      write(temp,"(a,'0',i2)")trim(name),myid
    else
      write(temp,"(a,'00',i1)")trim(name),myid
    endif
  endif
  unique=temp
  return
end function unique
```

Suggested exercises

Study, compile and run the program st_1d.f on various numbers of processors

Change it to use 2 or 1 MPI_bcast calls instead of 8

Hint: (The "correct" way to do it with 1 call is to use F90 and MPI derived data types)

Do the decomposition in rows

Do periodic boundary conditions

Modify the write_grid routine to output the whole grid from node 0

2d decomposition

The program is almost identical

We now have our grid distributed in a block fashion across the processors instead of striped

We can have ghost cells on 1, 2, 3 or 4 sides of the grid held on a particular processor

Example 2d Decomposition

50 x 50 grid on 4 processors


Grid on each processor is allocated to:

pid= 0 ($0 \leq i \leq 26$), ($0 \leq j \leq 26$)
pid= 1 ($0 \leq i \leq 26$), ($25 \leq j \leq 51$)
pid= 2 ($25 \leq i \leq 51$), ($0 \leq j \leq 26$)
pid= 3 ($25 \leq i \leq 51$), ($25 \leq j \leq 51$)

But each processor calculates only for:

pid= 0 ($1 \leq i \leq 25$), ($1 \leq j \leq 25$)
pid= 1 ($1 \leq i \leq 25$), ($26 \leq j \leq 50$)
pid= 2 ($26 \leq i \leq 50$), ($1 \leq j \leq 25$)
pid= 3 ($26 \leq i \leq 50$), ($26 \leq j \leq 50$)

Extra cells are ghost cells



**Grid Distributed across
4 processors**

Only three changes need to be made to our program

Given an arbitrary number of processors find a good topology (number of rows and columns of processors)

Make new communicators to allow for easy exchange of ghost cells

Set up communicators so that every processor in the same row is in a given communicator

Set up communicators so that every processor in the same column is in a given communicator

Add the up/down communication

Given an arbitrary number of processors find a good topology (number of rows and columns of processors)

```
nrow=nint(sqrt(float(numnodes)))
ncol=numnodes/nrow
do while (nrow*ncol .ne. numnodes)
    nrow=nrow+1
    ncol=numnodes/nrow
enddo
if(nrow .gt. ncol)then
    i=ncol
    ncol=nrow
    nrow=i
endif
myrow=myid/ncol+1
mycol=myid - (myrow-1)*ncol + 1
```

	nodes	nrow	ncol
	2	1	2
	3	3	1
	4	2	2
	5	5	1
	6	2	3
	7	7	1
	8	4	2
	9	3	3
	10	5	2
	11	11	1
	12	3	4
	13	13	1
	14	7	2
	15	5	3
	16	4	4

Make new communicators to allow for easy exchange of ghost cells

! make the row and col communicators

! all processors with the same row will be in the same ROW_COMM
call MPI_COMM_SPLIT(MPI_COMM_WORLD,myrow,mycol,ROW_COMM,mpi_err)
call MPI_COMM_RANK(ROW_COMM, myid_row, mpi_err)
call MPI_COMM_SIZE(ROW_COMM, nodes_row, mpi_err)

! all processors with the same col will be in the same COL_COMM
call MPI_COMM_SPLIT(MPI_COMM_WORLD,mycol,myrow,COL_COMM,mpi_err)
call MPI_COMM_RANK(COL_COMM, myid_col, mpi_err)
call MPI_COMM_SIZE(COL_COMM, nodes_col, mpi_err)

! find id of neighbors using the communicators created above
mytop =myid_col-1;if(mytop .lt. 0)mytop =MPI_PROC_NULL
mybot =myid_col+1;if(mybot .eq. nodes_col)mybot =MPI_PROC_NULL
myleft =myid_row-1;if(myleft .lt. 0)myleft =MPI_PROC_NULL
myright=myid_row+1;if(myright .eq. nodes_row)myright=MPI_PROC_NULL

Communication up/down

```
if(even(myid_row))then
! send to top
    call MPI_SEND(psi(i1,:),num_y,MPI_DOUBLE_PRECISION,mytop,      &
                  10, COL_COMM,mpi_err)
! rec from top
    call MPI_RECV(psi(i1-1,:),num_y,MPI_DOUBLE_PRECISION,mytop,    &
                  10,COL_COMM,status,mpi_err)
! rec from bot
    call MPI_RECV(psi(i2+1,:),num_y,MPI_DOUBLE_PRECISION,mybot,    &
                  10,COL_COMM,status,mpi_err)
! send to bot
    call MPI_SEND(psi(i2,:),num_y,MPI_DOUBLE_PRECISION,mybot,      &
                  10, COL_COMM,mpi_err)
else
```

Communication up/down (continued)

```
! rec from bot
    call MPI_RECV(psi(i2+1,:),num_y,MPI_DOUBLE_PRECISION,mybot,    &
                  10,COL_COMM,status,mpi_err)
! send to bot
    call MPI_SEND(psi(i2,:),num_y,MPI_DOUBLE_PRECISION,mybot,    &
                  10,COL_COMM,mpi_err)
! send to top
    call MPI_SEND(psi(i1,:),num_y,MPI_DOUBLE_PRECISION,mytop,    &
                  10,COL_COMM,mpi_err)
! rec from top
    call MPI_RECV(psi(i1-1,:),num_y,MPI_DOUBLE_PRECISION,mytop,    &
                  10,COL_COMM,status,mpi_err)
endif
```

We may want the master to do all output

Each processor holds a section of the grid

If each node writes we end up with a lot of pieces

If we have good communication it may be faster for 1 processor to write the file

We want a scalable solution

Write the grid from the master

Write a line (row) at a time

Do $l=i0,i3$ (A four step process)

- (1) Processors find how much of the line they hold
- (2) Processors tell master how much of the line they have (MPI_GATHER)
- (3) Processors send the data to the master (MPI_GATHERV)
- (4) Master prints the line

Enddo

Processors find how much of the line they hold

! Check to see if the present line is within
! the section that I hold

```
if(i .ge. istart .and. i .le. iend)then
```

```
    dj=jend-jstart+1
```

```
    mystart=jstart
```

```
    myend=jend
```

```
    icol=i
```

```
else ! I don't send anything
```

```
    dj=0
```

```
    mystart=jstart
```

```
    myend=jstart
```

```
    icol=istart
```

```
endif
```

Processors tell master how much of the line they have (MPI_GATHER)

```
call MPI_GATHER(dj      , 1, MPI_INTEGER, &  
               counts, 1, MPI_INTEGER, &  
               mpi_master,MPI_COMM_WORLD,mpi_err)
```

Each processor sends its "number to be sent" value in dj, type MPI_INTEGER

Mpi_master gets a "number to be sent" from each processor in the array counts

Processors send the data to the master (MPI_GATHERV)

```
if(myid .eq. mpi_master)then
  do k=1,numnodes-1
    offsets(k)=counts(k-1)+offsets(k-1)
  enddo
endif
call MPI_GATHERV(psi(icol,mystart:myend),dj,&
                 MPI_DOUBLE_PRECISION,      &
                 arow,counts,offsets,        &
                 MPI_DOUBLE_PRECISION,      &
                 mpi_master,MPI_COMM_WORLD,  &
                 mpi_err)
```

Processors send the data to the master (MPI_GATHERV)

Psi(icol,mystart:myend) : data being sent

Dj : number of values being sent

Arow : data ends up here

Counts : array that holds the number of values being sent from each processor

Offsets : array that holds a pointer to the beginning of values received from each processor, calculated from counts

Master writes the data

```
if(myid .eq. mpi_master)then
  do j=j0,j3
    write(18,'(g14.7)',advance="no")arow(j)
  enddo
  write(18,*)
endif
```

The answer is: (after 75,000 iterations)

200 x 200 interior grid size

2000000 x 2000000 physical size

$\alpha=1.0\text{e-}9$ $\beta=2.25\text{e-}11$ $\gamma=3.0\text{e-}6$

