

MPI Case Study Jacobi-Solver (Point-to-point communication)

Dieter an Mey

e-mail: anmey@rz.rwth-aachen.de

Jacobi-Algorithmus

Finite Difference Method to solve the 2D **Helmholtz equation**

$$\Delta u + \kappa u = f$$

with homogeneous Dirichlet-boundary conditions is solved with a **finite difference method**.

The Laplace operator Δ is discretized with the central **5-point difference star** .

The corresponding **lineare equations** with a banded coefficient matrix are solved iteratively with the (simple) **Jacobi method**.

The Jacobi-method can easily be **parallelized**.

This example is from the OpenMP web page www.openmp.org

```
#define U(i,j) u[(i)*n+(j)]
#define UOLD(i,j) uold[(i)*n+(j)]
#define F(i,j) f[(i)*n+(j)]


/* ... */
error = 10.0 * tol;
k = 1;
while (k <= maxit && error > tol) {
    error = 0.0;

    for (j=0; j<m; j++)
        for (i=0; i<n; i++)
            UOLD(j,i) = U(j,i);

    for (j=1; j<m-1; j++)
        for (i=1; i<n-1; i++){
            resid=(ax*(UOLD(j,i-1)+UOLD(j,i+1))+ay*(UOLD(j1,i)+UOLD(j+1,i))
                + b * UOLD(j,i) - F(j,i) ) / b;
            U(j,i) = UOLD(j,i) - omega * resid;
            error =error + resid*resid;
        }

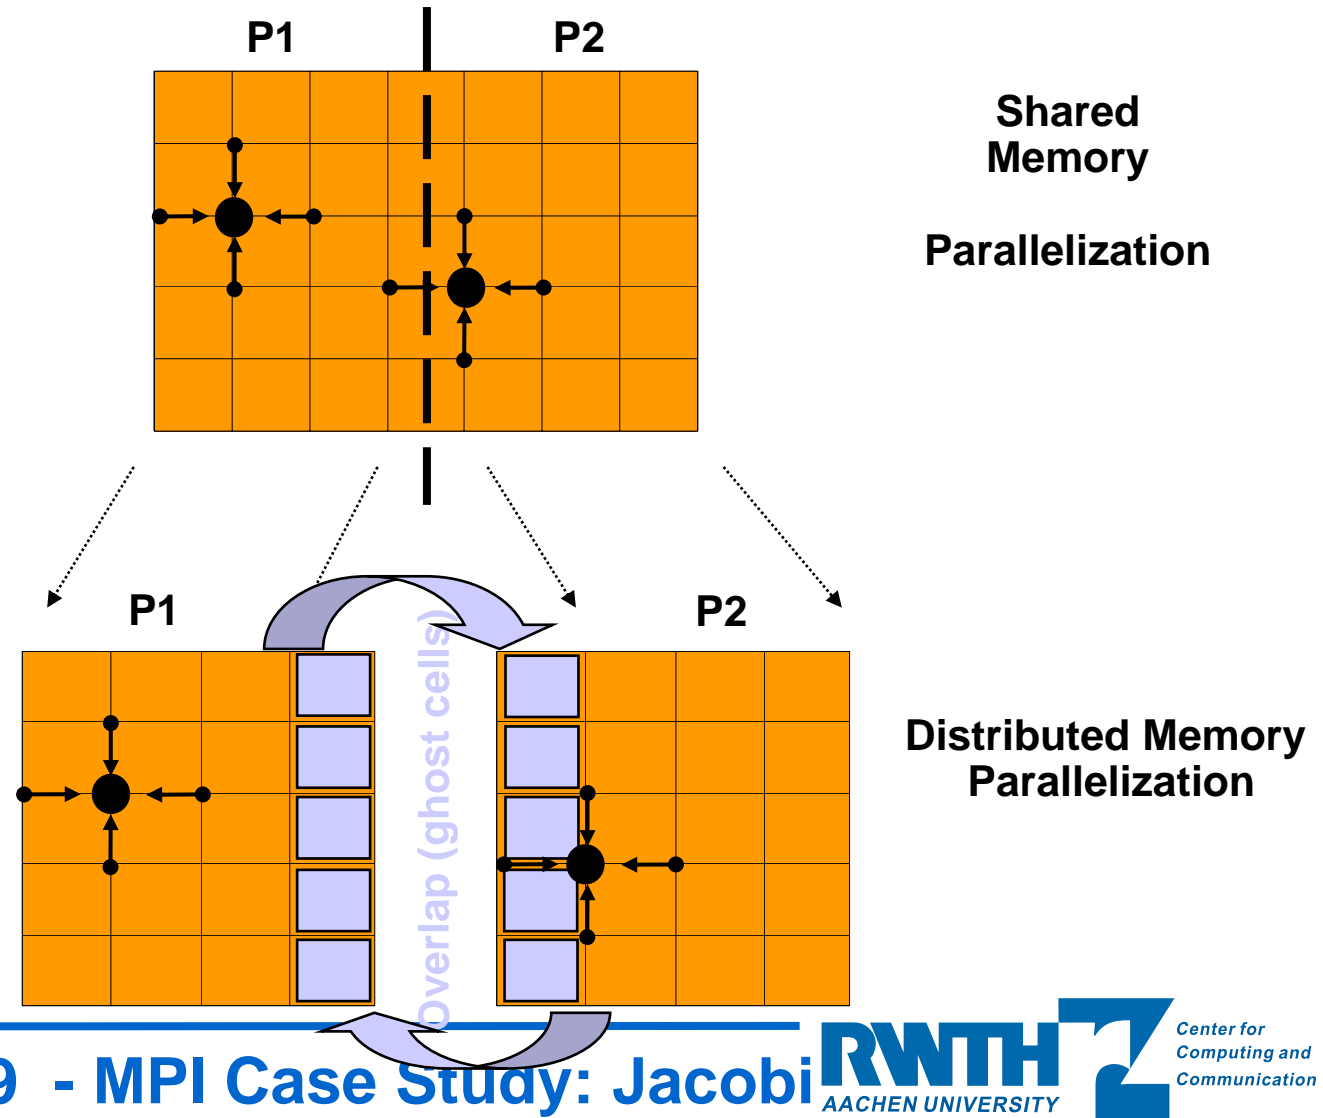
    k++;
    error = sqrt(error) /(n*m);
} /* while */
/* ... */
```

Jacobi-Method – Domain Decomposition



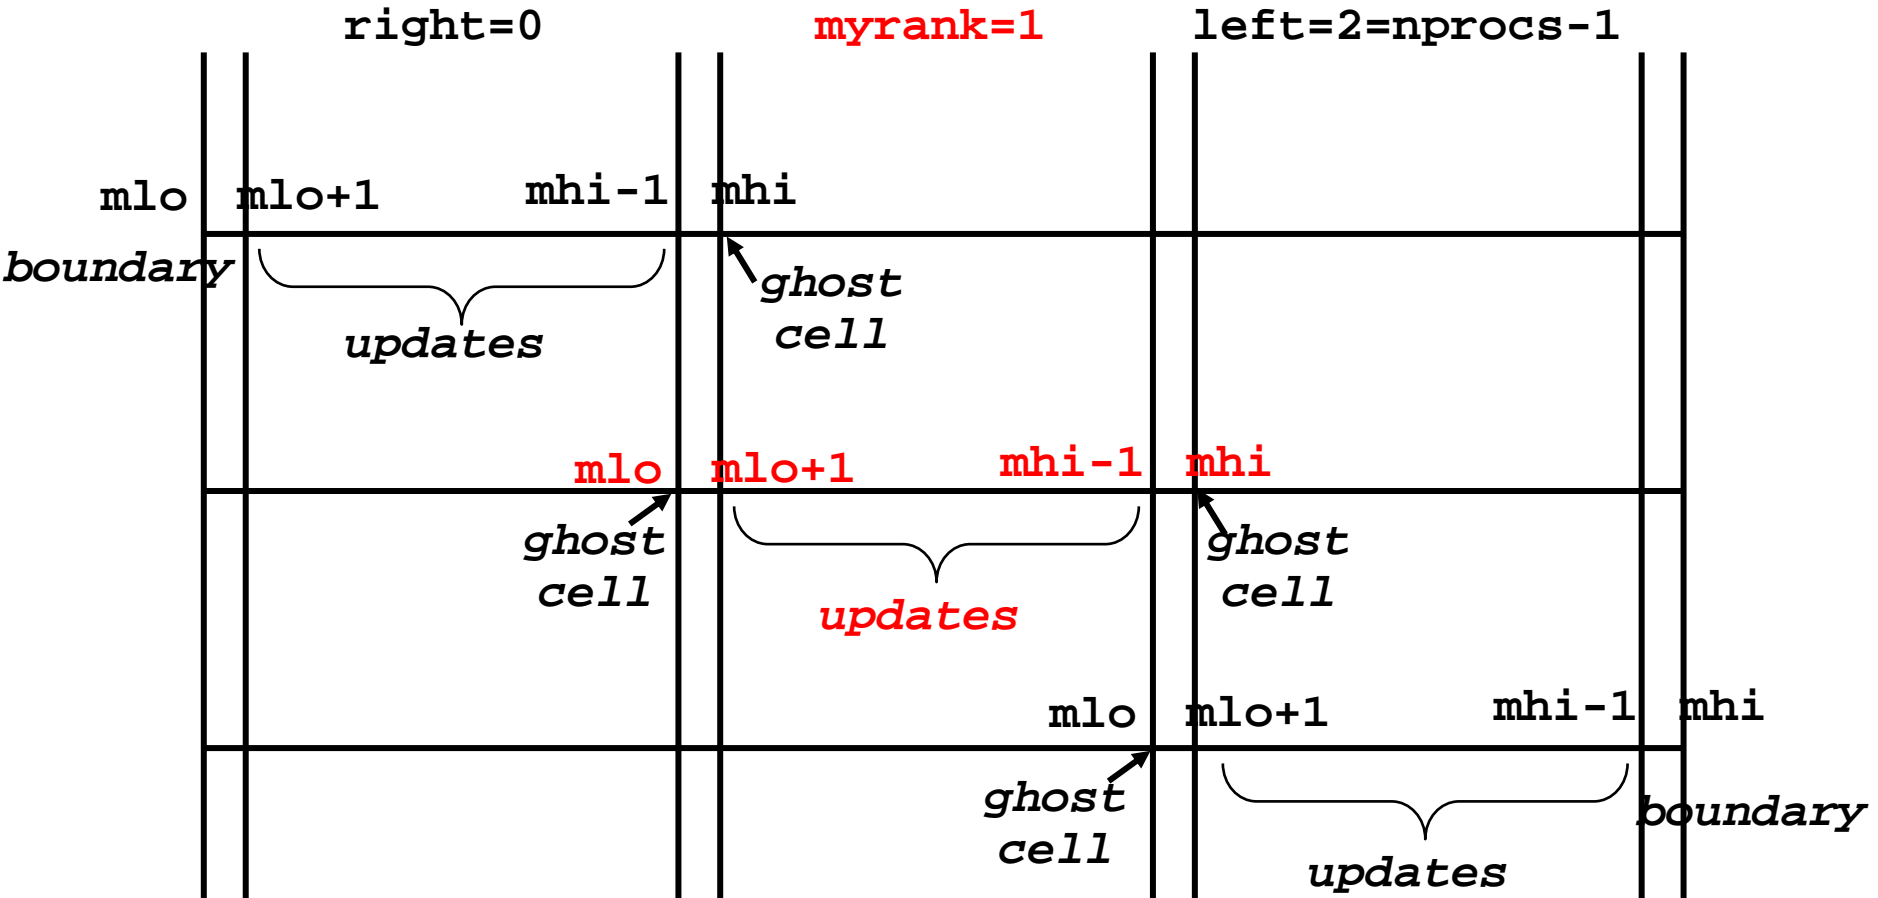
Shared
Memory
Parallelization

Jacobi-Method - Domain Decomposition



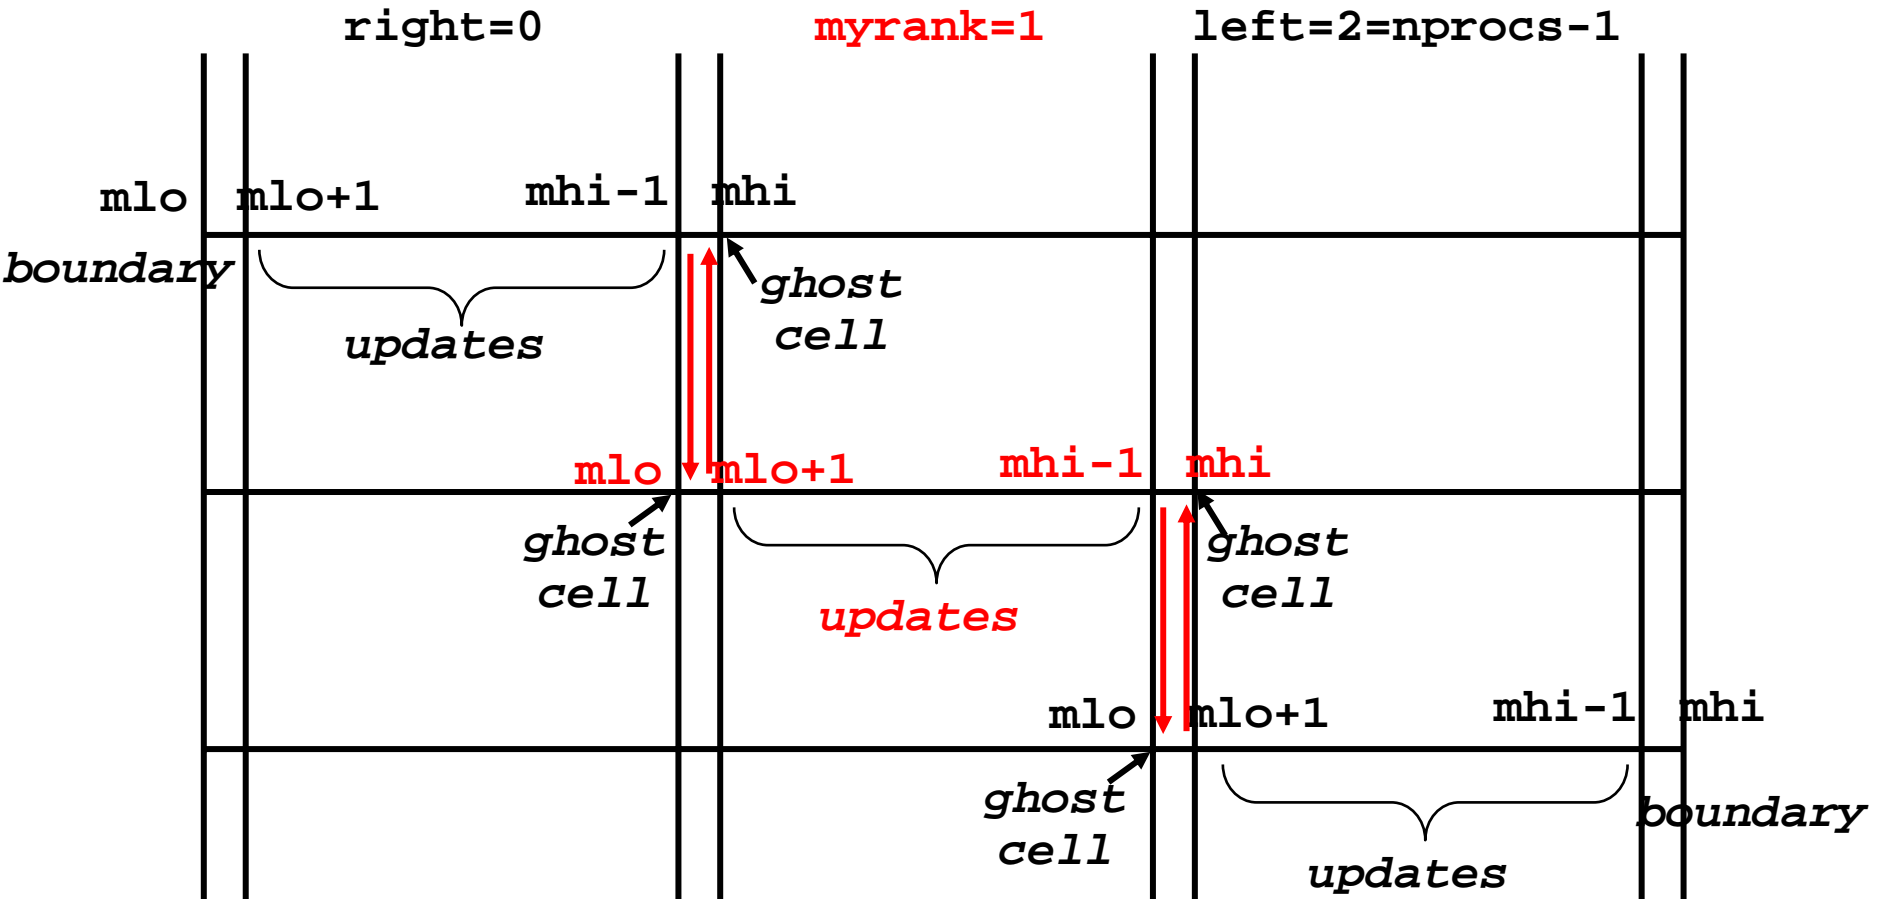
Jacobi-Method - Domain Decomposition

For example: $nprocs = 3$



Jacobi-Method - Domain Decomposition

For example: $nprocs = 3$





Jacobi – MPI Version 1: (B)locking send / recv ?

```
#define U(j,i) u[((j)-mlo)*n+(i)]
#define F(j,i) f[((j)-mlo)*n+(i)]
#define UOLD(j,i) uold[((j)-mlo)*n+(i)]

while (k <= maxit && error > tol) {
    error = 0.0;
    /* copy new solution into old - This is a very dangerous implementation. Why? */
    if (myrank != nprocs-1){ /* send stripe mhi-1 to right neighbour */
        MPI_Send(&U(mhi-1,0), n, MPI_DOUBLE, myrank+1,tag, MPI_COMM_WORLD);
    }
    if (myrank != 0){ /* send stripe mlo+1 to left neighbour */
        MPI_Send(&U(mlo+1,0), n, MPI_DOUBLE, myrank-1,tag, MPI_COMM_WORLD);
    }


    if (myrank != 0){ /* receive stripe mlo from left neighbour */
        MPI_Recv(&UOLD(mlo,0),n,MPI_DOUBLE,myrank-1,tag,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
    }
    if (myrank != nprocs-1 ){ /* receive stripe mhi from right neighbour */
        MPI_Recv(&UOLD(mhi,0),n,MPI_DOUBLE,myrank+1,tag,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
    }

    for (j=mlo+1; j<=mhi-1; j++) for (i=0; i<n; i++) UOLD(j,i) = U(j,i);
    for (j=mlo+1; j<=mhi-1; j++) for (i=1; i<n-1; i++){
        resid = ...;
        U(j,i) = UOLD(j,i) - omega * resid;
        error = error + resid*resid;
    }
    error_local = error;
    MPI_Allreduce(&error_local, &error, 1, MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);

    k++; error = sqrt(error) /(n*m);
} /* while */
```


Deadlock?

How about this? Will it deadlock?



MPI_Sendrecv

```
MPI_Sendrecv(void *sendbuf,
             int sendcount,
             MPI_Datatype sendtype,
             int dest,
             int sendtag,
             void *recvbuf,
             int recvcount,
             MPI_Datatype recvtype,
             int source,
             int recvtag,
             MPI_Comm comm,
             MPI_Status *status );
```

What?

How much?

Which type?
destination?

Message tag

To where ?

How much?

Which type?

From where?

Message tag

Communicator?

Receive Status?



Jacobi – MPI Version 2: Sendrecv

```
while (k <= maxit && error > tol) {
    error = 0.0;

    /* copy new solution into old */

    left = myrank-1;    if ( myrank == 0 )           left = MPI_PROC_NULL;
    right = myrank+1;   if ( myrank == nprocs-1 )     right = MPI_PROC_NULL;

    /* exchange stripe with right neighbour */
    MPI_Sendrecv(&U(mhi-1,0), n, MPI_DOUBLE, right, TAG_MOVE,
                &UOLD(mhi,0), n, MPI_DOUBLE, right, TAG_MOVE,
                MPI_COMM_WORLD, MPI_STATUS_IGNORE);

    /* exchange stripe with left neighbour */
    MPI_Sendrecv(&U(mlo+1,0), n, MPI_DOUBLE, left, TAG_MOVE,
                &UOLD(mlo,0), n, MPI_DOUBLE, left, TAG_MOVE,
                MPI_COMM_WORLD, MPI_STATUS_IGNORE);


    for (j=mlo+1; j<=mhi-1; j++) for (i=0; i<n; i++) UOLD(j,i) = U(j,i);
    for (j=mlo+1; j<=mhi-1; j++) for (i=1; i<n-1; i++){
        resid =(
            ax * (UOLD(j,i-1) + UOLD(j,i+1))
            + ay * (UOLD(j-1,i) + UOLD(j+1,i))
            + b * UOLD(j,i) - F(j,i)
        ) / b;
        U(j,i) = UOLD(j,i) - omega * resid;
        error = error + resid*resid;
    }
    error_local = error;
    MPI_Allreduce(&error_local, &error, 1, MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);
    k++; error = sqrt(error) /(n*m);
} /* while */
```



Jacobi – MPI Version 3: Asynchronous Send/Recv

```
MPI_Request request[4];  
MPI_Status status[4];  
int reqcnt, left, right;
```

```
/* copy new solution into old */  
left = myrank-1; if ( myrank == 0 ) left = MPI_PROC_NULL;  
right = myrank+1; if ( myrank == nprocs-1 ) right = MPI_PROC_NULL;  
  
reqcnt = 0;  
/* receive stripe mlo from left neighbour blocking */  
MPI_Irecv(&UOLD(mlo,0), n, MPI_DOUBLE, left,  
          TAG_MOVE_RIGHT, MPI_COMM_WORLD, &request[reqcnt]);  
reqcnt++;  
/* receive stripe mhi from right neighbour blocking */  
MPI_Irecv(&UOLD(mhi,0), n, MPI_DOUBLE, right,  
          TAG_MOVE_LEFT, MPI_COMM_WORLD, &request[reqcnt]);  
reqcnt++;  
/* send stripe mhi-1 to right neighbour async */  
MPI_Isend(&U(mhi-1,0), n, MPI_DOUBLE, right,  
          TAG_MOVE_RIGHT, MPI_COMM_WORLD, &request[reqcnt]);  
reqcnt++;  
/* send stripe mlo+1 to left neighbour async */  
MPI_Isend(&U(mlo+1,0), n, MPI_DOUBLE, left,  
          TAG_MOVE_LEFT, MPI_COMM_WORLD, &request[reqcnt]);  
  
for (j=mlo+1; j<=mhi-1; j++)  
    for (i=0; i<n; i++)  
        UOLD(j,i) = U(j,i);  
  
MPI_Waitall(reqcnt, request, status);
```



Overlap
Communication and
Computation



Jacobi Solver – OpenMP Version

```
#define U(i,j) u[(i)*n+(j)]
#define UOLD(i,j) uold[(i)*n+(j)]
#define F(i,j) f[(i)*n+(j)]
#include <omp.h>
/* ... */
    error = 10.0 * tol;
    k = 1;
    while (k <= maxit && error > tol) {
        error = 0.0;
#pragma omp parallel private(j,i,resid) \
                        shared(m,n,ax,ay,u,uold,f,b,omega)
{
#pragma omp for
    for (j=0; j<m; j++)
        for (i=0; i<n; i++)
            UOLD(j,i) = U(j,i);
#pragma omp for reduction(+:error)
    for (j=1; j<m-1; j++)
        for (i=1; i<n-1; i++){
            resid=(ax*(UOLD(j,i-1)+UOLD(j,i+1))+ay*(UOLD(j-1,i)+UOLD(j+1,i))
                + b * UOLD(j,i) - F(j,i) ) / b;
            U(j,i) = UOLD(j,i) - omega * resid;
            error =error + resid*resid;
        }
} /* end of parallel region */
    k++;
    error = sqrt(error) /(n*m);
} /* while */
/* ... */
```



Jacobi – MPI+OpenMP=Hybrid

```
while (k <= maxit && error > tol) {
    error = 0.0;

    /* copy new solution into old */

    left = myrank-1;    if ( myrank == 0 )          left = MPI_PROC_NULL;
    right = myrank+1;   if ( myrank == nprocs-1 )    right = MPI_PROC_NULL;

    /* exchange stripe with right neighbour */
    MPI_Sendrecv(&U(mhi-1,0), n, MPI_DOUBLE, right, TAG_MOVE,
                &UOLD(mhi,0), n, MPI_DOUBLE, right, TAG_MOVE,
                MPI_COMM_WORLD, MPI_STATUS_IGNORE);

    /* exchange stripe with left neighbour */
    MPI_Sendrecv(&U(mlo+1,0), n, MPI_DOUBLE, left, TAG_MOVE,
                &UOLD(mlo,0), n, MPI_DOUBLE, left, TAG_MOVE,
                MPI_COMM_WORLD, MPI_STATUS_IGNORE);

#pragma omp parallel for
    for (j=mlo+1; j<=mhi-1; j++) for (i=0; i<n; i++) UOLD(j,i) = U(j,i);
#pragma omp parallel for private(resid) reduction(+:error)
    for (j=mlo+1; j<=mhi-1; j++) for (i=1; i<n-1; i++){
        resid = (
            ax * (UOLD(j,i-1) + UOLD(j,i+1))
            + ay * (UOLD(j-1,i) + UOLD(j+1,i))
            + b * UOLD(j,i) - F(j,i)
        ) / b;
        U(j,i) = UOLD(j,i) - omega * resid;
        error = error + resid*resid;
    }
    error_local = error;
    MPI_Allreduce(&error_local, &error, 1, MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);
    k++; error = sqrt(error) / (n*m);
} /* while */
```